

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\fiat500_VehicleSelection_Dataset (1).csv")
a
```

Out[2]:

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | |
|------|-----|--------|--------------|-------------|----------|-----------------|-----------|--------|
| 0 | 1.0 | lounge | 51.0 | 882.0 | 25000.0 | 1.0 | 44.907242 | 8.6115 |
| 1 | 2.0 | pop | 51.0 | 1186.0 | 32500.0 | 1.0 | 45.666359 | 12.241 |
| 2 | 3.0 | sport | 74.0 | 4658.0 | 142228.0 | 1.0 | 45.503300 | 11 |
| 3 | 4.0 | lounge | 51.0 | 2739.0 | 160000.0 | 1.0 | 40.633171 | 17.634 |
| 4 | 5.0 | pop | 73.0 | 3074.0 | 106880.0 | 1.0 | 41.903221 | 12.495 |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 1544 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 1545 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 1546 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | Null |
| 1547 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 1548 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |

1549 rows × 11 columns



In [3]:

```
b=a.head(100)
b
```

Out[3]:

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon |
|-----|-------|--------|--------------|-------------|----------|-----------------|-----------|---------|
| 0 | 1.0 | lounge | 51.0 | 882.0 | 25000.0 | 1.0 | 44.907242 | 8.61155 |
| 1 | 2.0 | pop | 51.0 | 1186.0 | 32500.0 | 1.0 | 45.666359 | 12.2418 |
| 2 | 3.0 | sport | 74.0 | 4658.0 | 142228.0 | 1.0 | 45.503300 | 11.4 |
| 3 | 4.0 | lounge | 51.0 | 2739.0 | 160000.0 | 1.0 | 40.633171 | 17.6346 |
| 4 | 5.0 | pop | 73.0 | 3074.0 | 106880.0 | 1.0 | 41.903221 | 12.4956 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 95 | 96.0 | sport | 51.0 | 4292.0 | 165600.0 | 1.0 | 44.715408 | 11.3083 |
| 96 | 97.0 | pop | 51.0 | 1066.0 | 28000.0 | 1.0 | 41.769051 | 12.6628 |
| 97 | 98.0 | sport | 51.0 | 2009.0 | 86000.0 | 2.0 | 40.633171 | 17.6346 |
| 98 | 99.0 | lounge | 51.0 | 456.0 | 18592.0 | 2.0 | 45.393600 | 10.4822 |
| 99 | 100.0 | pop | 51.0 | 731.0 | 41558.0 | 2.0 | 45.571220 | 9.15913 |

100 rows × 11 columns

In [4]:

```
b.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 11 columns):
#   Column              Non-Null Count  Dtype
---  -
0   ID                   100 non-null   float64
1   model                100 non-null   object
2   engine_power         100 non-null   float64
3   age_in_days          100 non-null   float64
4   km                   100 non-null   float64
5   previous_owners      100 non-null   float64
6   lat                  100 non-null   float64
7   lon                  100 non-null   object
8   price                100 non-null   object
9   Unnamed: 9           0 non-null     float64
10  Unnamed: 10          0 non-null     object
dtypes: float64(7), object(4)
memory usage: 8.7+ KB
```

In [5]:

```
b.describe()
```

Out[5]:

| | ID | engine_power | age_in_days | km | previous_owners | lat |
|-------|------------|--------------|-------------|---------------|-----------------|------------|
| count | 100.000000 | 100.000000 | 100.000000 | 100.000000 | 100.000000 | 100.000000 |
| mean | 50.500000 | 53.010000 | 1935.300000 | 58812.180000 | 1.180000 | 43.612648 |
| std | 29.011492 | 6.014284 | 1414.251278 | 44728.034639 | 0.500101 | 2.083451 |
| min | 1.000000 | 51.000000 | 366.000000 | 4000.000000 | 1.000000 | 38.218128 |
| 25% | 25.750000 | 51.000000 | 723.500000 | 19781.750000 | 1.000000 | 41.744165 |
| 50% | 50.500000 | 51.000000 | 1446.000000 | 44032.000000 | 1.000000 | 44.831066 |
| 75% | 75.250000 | 51.000000 | 3265.500000 | 95075.750000 | 1.000000 | 45.396568 |
| max | 100.000000 | 74.000000 | 4658.000000 | 188000.000000 | 3.000000 | 46.176498 |

In [6]:

```
c=b.dropna(axis=1)
c
```

Out[6]:

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | |
|-----|-------|--------|--------------|-------------|----------|-----------------|-----------|---------|
| 0 | 1.0 | lounge | 51.0 | 882.0 | 25000.0 | 1.0 | 44.907242 | 8.61155 |
| 1 | 2.0 | pop | 51.0 | 1186.0 | 32500.0 | 1.0 | 45.666359 | 12.2418 |
| 2 | 3.0 | sport | 74.0 | 4658.0 | 142228.0 | 1.0 | 45.503300 | 11.4 |
| 3 | 4.0 | lounge | 51.0 | 2739.0 | 160000.0 | 1.0 | 40.633171 | 17.6346 |
| 4 | 5.0 | pop | 73.0 | 3074.0 | 106880.0 | 1.0 | 41.903221 | 12.4956 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 95 | 96.0 | sport | 51.0 | 4292.0 | 165600.0 | 1.0 | 44.715408 | 11.3083 |
| 96 | 97.0 | pop | 51.0 | 1066.0 | 28000.0 | 1.0 | 41.769051 | 12.6628 |
| 97 | 98.0 | sport | 51.0 | 2009.0 | 86000.0 | 2.0 | 40.633171 | 17.6346 |
| 98 | 99.0 | lounge | 51.0 | 456.0 | 18592.0 | 2.0 | 45.393600 | 10.4822 |
| 99 | 100.0 | pop | 51.0 | 731.0 | 41558.0 | 2.0 | 45.571220 | 9.15913 |

100 rows × 9 columns

In [7]:

```
c.columns
```

Out[7]:

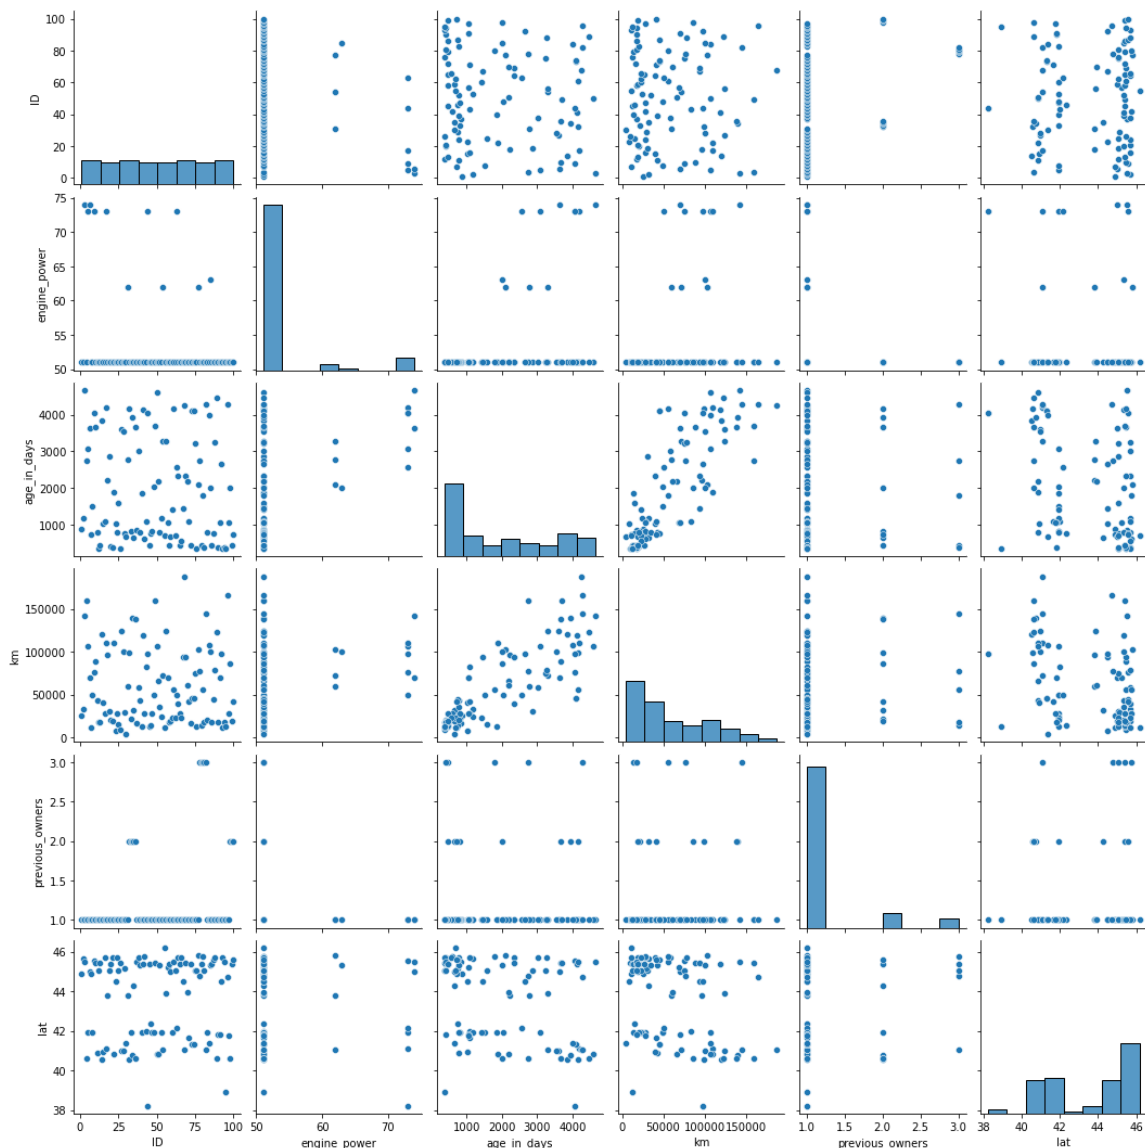
```
Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
      'lat', 'lon', 'price'],  
      dtype='object')
```

In [8]:

```
sns.pairplot(c)
```

Out[8]:

<seaborn.axisgrid.PairGrid at 0x27032bf0c70>



In [9]:


```
sns.distplot(c['age_in_days'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[9]:

<AxesSubplot:xlabel='age_in_days', ylabel='Density'>



In [10]:


```
f=c[['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
    'lat', 'lon', 'price']]
```

In [11]:

```
sns.heatmap(f.corr())
```

Out[11]:

<AxesSubplot:>



In [12]:

```
x=f[['ID','engine_power','km', 'previous_owners',  
      'lat', 'lon', 'price']]  
y=f['age_in_days']
```

In [13]:

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.5)
```

In [14]:

```
from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[14]:

LinearRegression()

In [15]:

```
print(lr.intercept_)
```

6896.681856091886

In [16]:

```
r=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
r
```

Out[16]:


| | Co-efficient |
|-----------------|--------------|
| ID | -3.021181 |
| engine_power | 17.469947 |
| km | -0.000558 |
| previous_owners | 165.049390 |
| lat | -13.430399 |
| lon | -21.678906 |
| price | -0.620086 |

In [17]:

```
u=lr.predict(x_test)  
plt.scatter(y_test,u)
```

Out[17]:

<matplotlib.collections.PathCollection at 0x2703521fb80>



In [18]:

```
print(lr.score(x_test,y_test))
```

0.8236052465600654

In [19]:

```
lr.score(x_train,y_train)
```

Out[19]:

0.8554274608268977

RIDGE REGRESSION

In [20]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [21]:

```
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[21]:

Ridge(alpha=10)

In [22]:

```
rr.score(x_test,y_test)
```

Out[22]:

0.8285017241960643

LASSO REGRESSION

In [23]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[23]:

Lasso(alpha=10)

In [24]:

```
la.score(x_test,y_test)
```

Out[24]:

0.8249223223241258

ELASTIC NET

In [25]:

```
from sklearn.linear_model import ElasticNet  
p=ElasticNet()  
p.fit(x_train,y_train)
```

Out[25]:

ElasticNet()

In [26]:

```
print(p.coef_)  
  
[-2.57150661e+00  1.64724226e+01 -2.26919036e-04  5.22363640e+01  
 -5.96119328e+00 -1.40467836e+01 -6.14410250e-01]
```

In [27]:

```
print(p.intercept_)  
  
6586.73754164688
```

In [28]:

```
print(p.predict(x_test))  
  
[ 534.69330542 3230.73561288 3899.69912776 2135.09156192 1045.28991093  
 640.36875073 3284.45983865 2281.55023196 3767.8667207  3379.48686212  
 966.71728847 4057.55919624 3321.43056837 2589.03223299  820.19949251  
2655.63509893 3578.3273058  2311.12233595 2082.57420273  407.96270267  
 866.62473865 2099.8057086   288.76544001  157.59228188  330.59010419  
 210.93087482 3963.07151446 3845.40545871 3256.08115958 1075.55784732  
 987.98095935 3721.75037752 2718.87031935  175.15443212  194.57449918  
2180.72530995 2193.44122312 3113.48160811 1463.74145497  535.74522519  
2230.93206031  267.70141967 2305.17155892  245.55587025  536.75110676  
 689.39157829 2289.77152188 2016.21723854 2335.00574079 2549.31018213]
```

In [35]:

```
prediction=p.predict(x_test)  
print(p.score(x_test,y_test))  
  
0.830116976100481
```

EVALUATION METRICS

In [36]:

```
from sklearn import metrics
```

In [37]:

```
print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))  
  
Mean Absolytre Error: 475.27204122558624
```

In [38]:

```
print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))  
  
Mean Squared Error: 354735.83604626363
```

In [39]:

```
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root Mean Squared Error: 595.5970416701746

In []: