

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

14.IRIS

In [2]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\14_Iris.csv")
a
```

Out[2]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...
145	146	6.7	3.0	5.2	2.3	Iris-virginica
146	147	6.3	2.5	5.0	1.9	Iris-virginica
147	148	6.5	3.0	5.2	2.0	Iris-virginica
148	149	6.2	3.4	5.4	2.3	Iris-virginica
149	150	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 6 columns

In [3]:

```
b=a.head(100)
b
```

Out[3]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...
95	96	5.7	3.0	4.2	1.2	Iris-versicolor
96	97	5.7	2.9	4.2	1.3	Iris-versicolor
97	98	6.2	2.9	4.3	1.3	Iris-versicolor
98	99	5.1	2.5	3.0	1.1	Iris-versicolor
99	100	5.7	2.8	4.1	1.3	Iris-versicolor

100 rows × 6 columns

In [4]:

```
b.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0    Id              100 non-null    int64
1    SepalLengthCm   100 non-null    float64
2    SepalWidthCm    100 non-null    float64
3    PetalLengthCm   100 non-null    float64
4    PetalWidthCm    100 non-null    float64
5    Species         100 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 4.8+ KB
```

In [5]:

```
b.describe()
```

Out[5]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	100.000000	100.000000	100.000000	100.000000	100.000000
mean	50.500000	5.471000	3.094000	2.862000	0.785000
std	29.011492	0.641698	0.476057	1.448565	0.566288
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	25.750000	5.000000	2.800000	1.500000	0.200000
50%	50.500000	5.400000	3.050000	2.450000	0.800000
75%	75.250000	5.900000	3.400000	4.325000	1.300000
max	100.000000	7.000000	4.400000	5.100000	1.800000

In [6]:

```
b.columns
```

Out[6]:

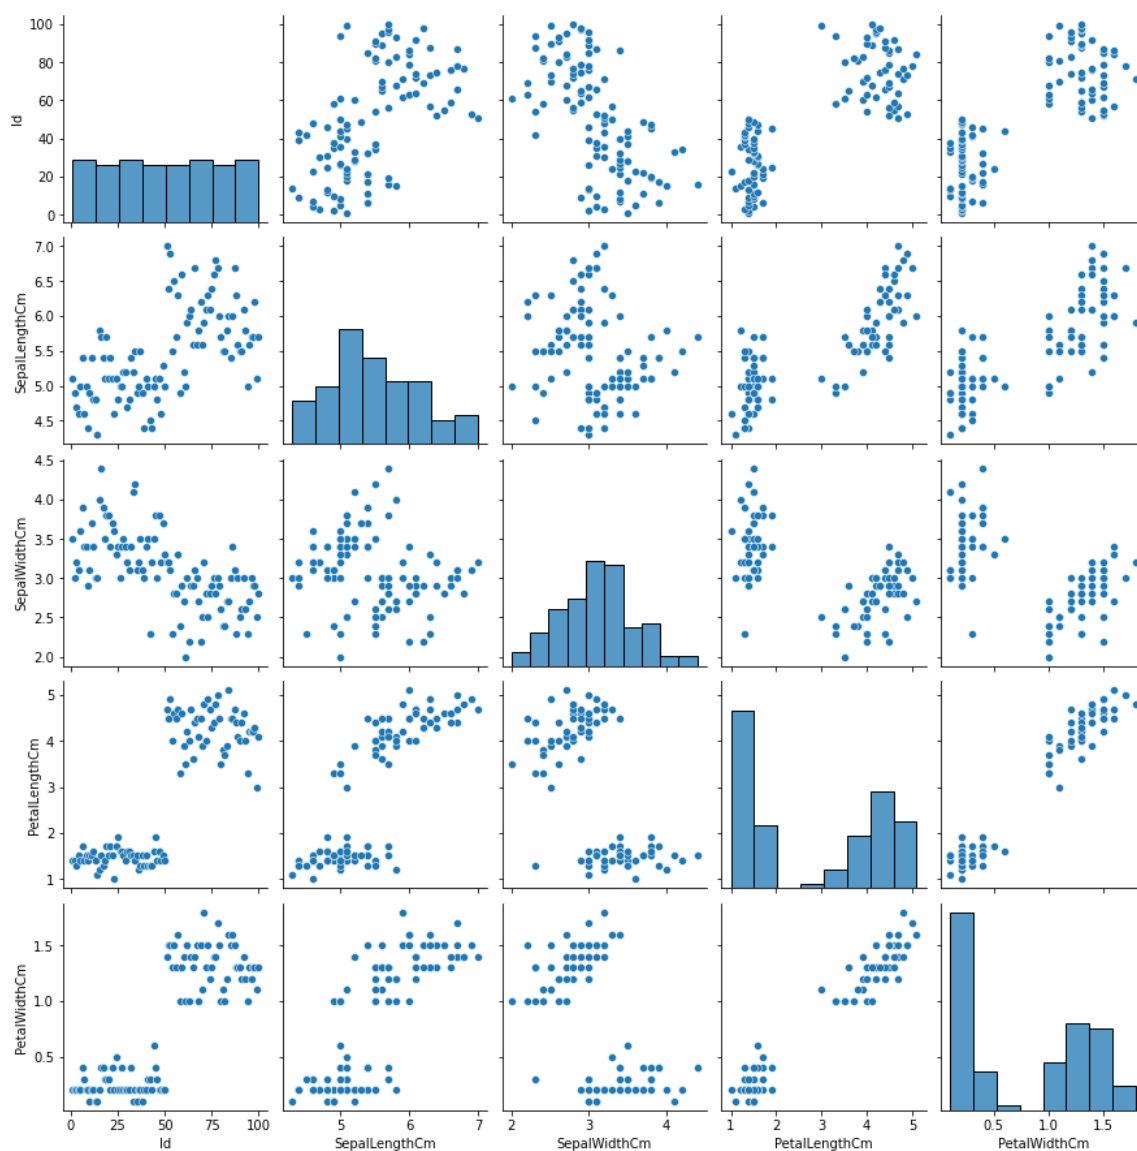
```
Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',  
      'Species'],  
      dtype='object')
```

In [7]:

```
sns.pairplot(b)
```

Out[7]:

<seaborn.axisgrid.PairGrid at 0x171d5960af0>



In [8]:


```
sns.distplot(b['PetalWidthCm'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[8]:

<AxesSubplot:xlabel='PetalWidthCm', ylabel='Density'>



In [9]:


```
f=b[['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',  
     'Species']]
```

In [10]:

```
sns.heatmap(f.corr())
```

Out[10]:

<AxesSubplot:>



In [11]:

```
x1=f['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm']  
y1=f['PetalWidthCm']
```

In [12]:

```
from sklearn.model_selection import train_test_split  
x1_train,x1_test,y1_train,y1_test=train_test_split(x1,y1,test_size=0.5)
```

In [13]:

```
from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()  
lr.fit(x1_train,y1_train)
```

Out[13]:

LinearRegression()

In [14]:

```
print(lr.intercept_)
```

-0.4226823928478569

In [15]:

```
r=pd.DataFrame(lr.coef_,x1.columns,columns=['Co-efficient'])  
r
```

Out[15]:


	Co-efficient
Id	0.001062
SepalLengthCm	-0.044434
SepalWidthCm	0.078301
PetalLengthCm	0.402064

In [16]:

```
u=lr.predict(x1_test)  
plt.scatter(y1_test,u)
```

Out[16]:

<matplotlib.collections.PathCollection at 0x171d79adc10>



In [17]:

```
print(lr.score(x1_test,y1_test))
```

0.9515948448055147

In [18]:

```
lr.score(x1_train,y1_train)
```

Out[18]:

0.9693146372994317

RIDGE REGRESSION

In [19]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [20]:

```
rr=Ridge(alpha=10)  
rr.fit(x1_train,y1_train)
```

Out[20]:

Ridge(alpha=10)

In [21]:

```
rr.score(x1_test,y1_test)
```

Out[21]:

0.9356412208675733

LASSO REGRESSION

In [22]:

```
la=Lasso(alpha=10)  
la.fit(x1_train,y1_train)
```

Out[22]:

Lasso(alpha=10)

In [23]:

```
la.score(x1_test,y1_test)
```

Out[23]:

0.23929107734430222

ELASTIC NET

In [25]:

```
from sklearn.linear_model import ElasticNet  
p=ElasticNet()  
p.fit(x1_train,y1_train)
```

Out[25]:

ElasticNet()

In [26]:

```
print(p.coef_)
```

[0.01540332 0. -0. 0.]

In [27]:

```
print(p.intercept_)
```

0.002953725623672354

In [28]:

```
print(p.predict(x1_test))
```

```
[1.46626882 1.51247877 0.81932951 0.37263333 1.09658921 1.32763896
 0.61908639 1.5432854 1.42005886 1.05037926 1.29683233 1.52788208
 0.11077694 1.14279916 1.28142901 0.92715273 1.37384891 0.26481011
 1.15820248 1.38925223 1.03497595 0.35723001 0.95795936 0.32642338
 1.01957263 1.40465555 0.18779353 0.20319684 0.52666649 0.91174941
 0.09537363 0.74231293 0.58827976 0.72690961 0.55747313 0.34182669
 0.7115063 0.69610298 0.24940679 0.40343996 0.23400348 1.06578258
 0.988766 1.11199253 0.28021343 0.51126318 1.43546218 1.12739585
 0.04916368 0.63448971]
```

In [33]:

```
prediction=p.predict(x1_test)
print(p.score(x1_test,y1_test))
```

0.6786581507031236

In [34]:

```
from sklearn import metrics
```

In [35]:

```
print("Mean Absolytre Error:",metrics.mean_absolute_error(y1_test,prediction))
```

Mean Absolytre Error: 0.27009296813033523

In [36]:

```
print("Mean Squared Error:",metrics.mean_squared_error(y1_test,prediction))
```

Mean Squared Error: 0.10417131533766419

In [37]:

```
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y1_test,prediction)))
```

Root Mean Squared Error: 0.32275581379374746

15.Horse Racing Results


In [38]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\15_Horse Racing Results.csv - 15_Horse Racing Res
a
```

Out[38]:

	Dato	Track	Race Number	Distance	Surface	Prize money	Starting position	Jockey	Jockey weight	C
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	6	K C Leung	52	1
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	14	C Y Ho	52	1
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	8	C Y Ho	52	1
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	13	Brett Prebble	54	1
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	9	C Y Ho	52	1
...
27003	14.06.2020	Sha Tin	11	1200	Gress	1450000	6	A Hamelin	59	A
27004	21.06.2020	Sha Tin	2	1200	Gress	967000	7	K C Leung	57	A
27005	21.06.2020	Sha Tin	4	1200	Gress	967000	6	Blake Shinn	57	A
27006	21.06.2020	Sha Tin	5	1200	Gress	967000	14	Joao Moreira	57	z
27007	21.06.2020	Sha Tin	11	1200	Gress	1450000	7	C Schofield	55	z

27008 rows × 21 columns




In [39]:

```
b=a.head(100)
b
```

Out[39]:

	Dato	Track	Race Number	Distance	Surface	Prize money	Starting position	Jockey	Jockey weight	Col
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	6	K C Leung	52	Sv
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	14	C Y Ho	52	Sv
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	8	C Y Ho	52	Sv
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	13	Brett Prebble	54	Sv
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	9	C Y Ho	52	Sv
...	
95	10.12.2017	Sha Tin	5	1200	Gress	18500000	13	Francois- Xavier Bertras	57	(B
96	10.12.2017	Sha Tin	7	1600	Gress	23000000	11	Ryan Moore	57	
97	01.10.2017	Sha Tin	7	1000	Gress	3000000	10	Brett Prebble	59	Zee
98	22.10.2017	Sha Tin	7	1200	Gress	4000000	9	Brett Prebble	59	Zee
99	19.11.2017	Sha Tin	7	1200	Gress	4000000	3	Brett Prebble	56	Zee

100 rows × 21 columns



In [40]:

b.describe()

Out[40]:

	Race Number	Distance	Prize money	Starting position	Jockey weight	Horse age	Path
count	100.000000	100.000000	1.000000e+02	100.000000	100.000000	100.000000	100.000000
mean	6.910000	1446.000000	3.562200e+06	6.170000	55.870000	6.580000	1.510000
std	2.099038	334.820923	4.486259e+06	3.440857	2.942736	1.35721	1.573101
min	1.000000	1000.000000	9.200000e+05	1.000000	49.000000	3.000000	0.000000
25%	6.000000	1200.000000	1.380000e+06	3.000000	54.000000	6.000000	0.000000
50%	7.000000	1400.000000	1.950000e+06	6.000000	56.000000	7.000000	1.000000
75%	8.000000	1650.000000	3.000000e+06	9.000000	58.000000	8.000000	3.000000
max	10.000000	2400.000000	2.300000e+07	14.000000	60.000000	9.000000	6.000000

In [41]:

b.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Dato                   100 non-null   object
1   Track                  100 non-null   object
2   Race Number           100 non-null   int64
3   Distance               100 non-null   int64
4   Surface                100 non-null   object
5   Prize money            100 non-null   int64
6   Starting position      100 non-null   int64
7   Jockey                 100 non-null   object
8   Jockey weight          100 non-null   int64
9   Country                100 non-null   object
10  Horse age              100 non-null   int64
11  TrainerName            100 non-null   object
12  Race time              100 non-null   object
13  Path                   100 non-null   int64
14  Final place            100 non-null   int64
15  FGrating               100 non-null   int64
16  Odds                   100 non-null   object
17  RaceType               100 non-null   object
18  HorseId                100 non-null   int64
19  JockeyId               100 non-null   int64
20  TrainerID              100 non-null   int64
dtypes: int64(12), object(9)
memory usage: 16.5+ KB
```

In [42]:

```
c=b.dropna(axis=1)
c
```

Out[42]:

	Dato	Track	Race Number	Distance	Surface	Prize money	Starting position	Jockey	Jockey weight	Col
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	6	K C Leung	52	Sv
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	14	C Y Ho	52	Sv
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	8	C Y Ho	52	Sv
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	13	Brett Prebble	54	Sv
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	9	C Y Ho	52	Sv
...
95	10.12.2017	Sha Tin	5	1200	Gress	18500000	13	Francois-Xavier Bertras	57	C B
96	10.12.2017	Sha Tin	7	1600	Gress	23000000	11	Ryan Moore	57	
97	01.10.2017	Sha Tin	7	1000	Gress	3000000	10	Brett Prebble	59	Zee
98	22.10.2017	Sha Tin	7	1200	Gress	4000000	9	Brett Prebble	59	Zee
99	19.11.2017	Sha Tin	7	1200	Gress	4000000	3	Brett Prebble	56	Zee

100 rows × 21 columns

In [43]:

```
c.columns
```

Out[43]:

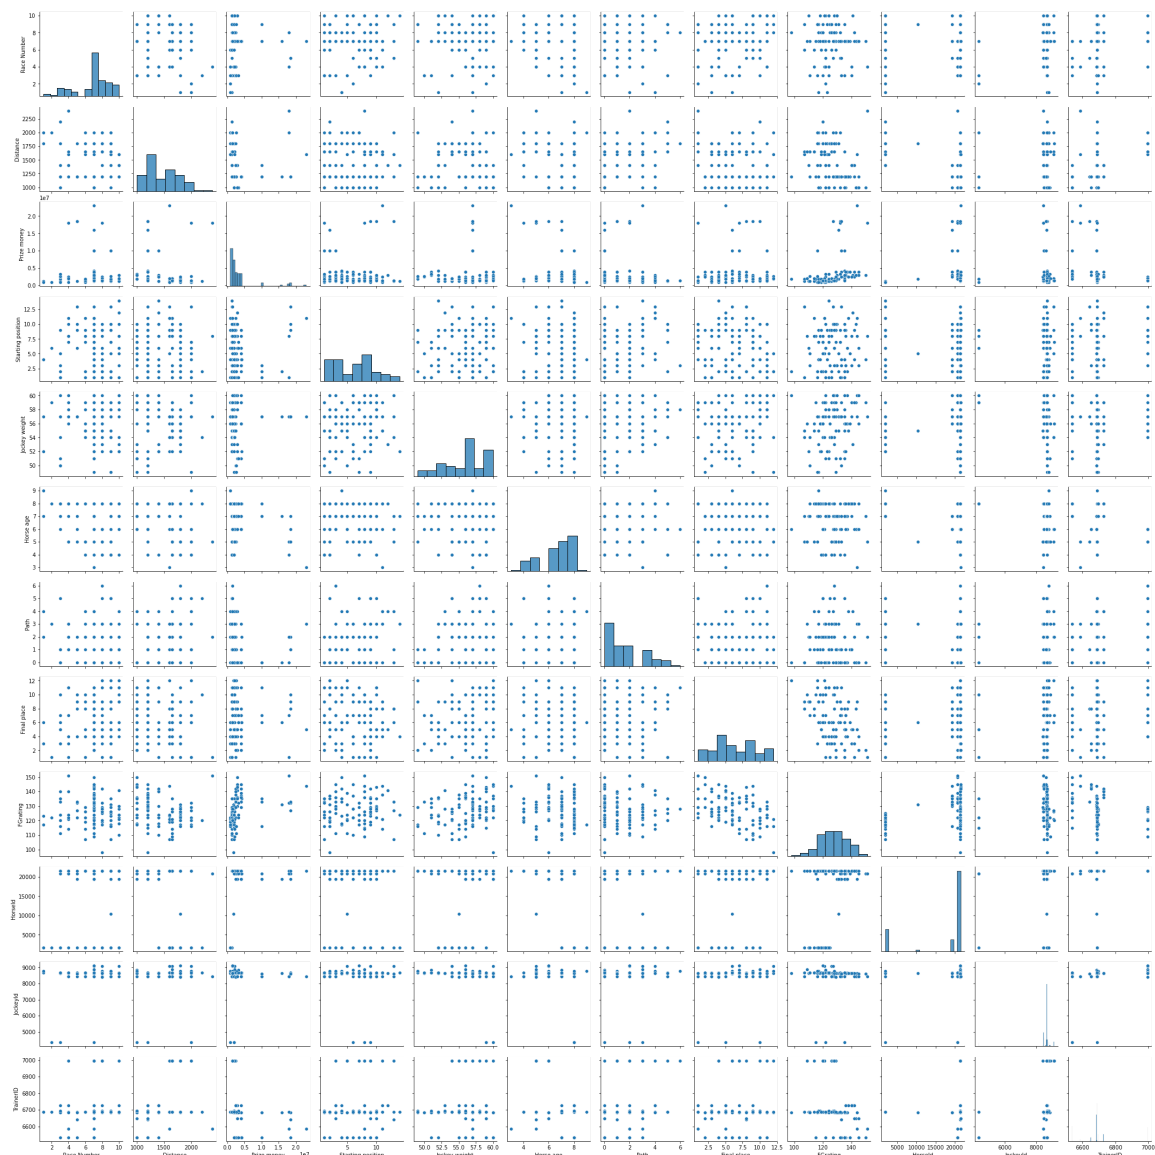
```
Index(['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',
      'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age',
      'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odd s',
      'RaceType', 'HorseId', 'JockeyId', 'TrainerID'],
      dtype='object')
```

In [44]:

sns.pairplot(c)

Out[44]:

<seaborn.axisgrid.PairGrid at 0x171d79f8160>



In [45]:


```
sns.distplot(c['Prize money'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[45]:

<AxesSubplot:xlabel='Prize money', ylabel='Density'>



In [46]:


```
f=c[['Dato', 'Race Number', 'Distance', 'Prize money',  
    'Starting position', 'Jockey weight', 'Horse age',  
    'Race time', 'Path', 'Final place', 'FGrating', 'Odds',  
    'HorseId', 'JockeyId', 'TrainerID']]
```

In [47]:

```
sns.heatmap(f.corr())
```

Out[47]:

<AxesSubplot:>



In [48]:

```
x2=f[['Distance','Starting position','Jockey weight','Horse age','FG rating','TrainerID']]
y2=f['Prize money']
```

In [49]:

```
from sklearn.model_selection import train_test_split
x2_train,x2_test,y2_train,y2_test=train_test_split(x2,y2,test_size=0.5)
```

In [50]:

```
from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x2_train,y2_train)
```

Out[50]:

LinearRegression()

In [51]:

```
print(lr.intercept_)
```

-7943479.703428032

In [52]:

```
r=pd.DataFrame(lr.coef_,x2.columns,columns=['Co-efficient'])
r
```

Out[52]:


	Co-efficient
Distance	-1405.295242
Starting position	-144265.340447
Jockey weight	16290.610897
Horse age	34287.629446
FGrating	28925.877574
TrainerID	1254.837090

In [53]:

```
u=lr.predict(x2_test)
plt.scatter(y2_test,u)
```

Out[53]:

<matplotlib.collections.PathCollection at 0x171dfb31a00>



In [54]:

```
print(lr.score(x2_test,y2_test))
```

-0.18099745594645333

In [55]:

```
lr.score(x2_train,y2_train)
```

Out[55]:

0.20270317432237261

In [56]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [57]:

```
rr=Ridge(alpha=10)  
rr.fit(x2_train,y2_train)
```

Out[57]:

Ridge(alpha=10)

In [58]:

```
rr.score(x2_test,y2_test)
```

Out[58]:

-0.17854010111414875

In [59]:

```
la=Lasso(alpha=10)  
la.fit(x2_train,y2_train)
```

Out[59]:

Lasso(alpha=10)

In [60]:

```
la.score(x2_test,y2_test)
```

Out[60]:

-0.18099496539455084

ELASTIC NET

In [61]:

```
from sklearn.linear_model import ElasticNet  
p=ElasticNet()  
p.fit(x2_train,y2_train)
```

Out[61]:

ElasticNet()

In [62]:

```
print(p.coef_)
```

```
[ -1412.0454908  -137098.97426212  14253.01782617  22168.26556704  
 28716.42411684   1161.8394489 ]
```

In [63]:

```
print(p.intercept_)
```

-7129709.166710248

In [64]:

```
print(p.predict(x2_test))
```

```
[2893874.22892909 3284561.6081243 2139951.88279886 2300236.85974332
 2066619.66203657 1400500.6780598 2012053.03819022 3006474.23646753
 3424680.98654762 2703709.07931783 1822122.98536334 1609180.80439627
 2523405.94488993 2369984.33402477 2680927.72586478 2785710.0442954
 2282522.67889327 1767733.98777093 3272606.42229179 2649741.03803308
 1514504.31946627 1260877.91388004 1768407.38197416 1650989.13849134
 3021463.12415376 2803876.65127175 3497843.46505108 2719847.68091414
 3050060.24664378 3585831.74535887 2251600.55928487 3396824.60664115
 2044543.45814441 1851873.11055457 1511915.51732091 1977888.39991734
 1730711.77762278 3035451.46411633 1156516.2672989 2785991.33833103
 2393346.81167774 2958553.4229894 1036481.90569666 1200433.27773343
 1610953.31873499 1737195.42596711 2443679.9453561 1294746.01786834
 3093444.57217237 1874294.40022923]
```

In [65]:

```
print(p.score(x2_test,y2_test))
```

-0.17561248978599076

EVALUATION METRICS

In [66]:

```
from sklearn import metrics
```

In [67]:

```
print("Mean Absolytre Error:",metrics.mean_absolute_error(y2_test,prediction))
```

Mean Absolytre Error: 4704399.176357559

In [68]:

```
print("Mean Squared Error:",metrics.mean_squared_error(y2_test,prediction))
```

Mean Squared Error: 56256648448046.99

In [69]:

```
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y2_test,prediction)))
```

Root Mean Squared Error: 7500443.216773725

16_Sleep_health_and_lifestyle_dataset


In [71]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\16_Sleep_health_and_lifestyle_dataset.csv")
a
```

Out[71]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Pr
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	
1	2	Male	28	Doctor	6.2	6	60	8	Normal	
2	3	Male	28	Doctor	6.2	6	60	8	Normal	
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	
...
369	370	Female	59	Nurse	8.1	9	75	3	Overweight	
370	371	Female	59	Nurse	8.0	9	75	3	Overweight	
371	372	Female	59	Nurse	8.1	9	75	3	Overweight	
372	373	Female	59	Nurse	8.1	9	75	3	Overweight	
373	374	Female	59	Nurse	8.1	9	75	3	Overweight	

374 rows × 13 columns



In [72]:

```
b=a.head(100)
b
```

Out[72]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Pre
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	1
1	2	Male	28	Doctor	6.2	6	60	8	Normal	1
2	3	Male	28	Doctor	6.2	6	60	8	Normal	1
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	1
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	1
...
95	96	Female	36	Accountant	7.1	8	60	4	Normal	1
96	97	Female	36	Accountant	7.2	8	60	4	Normal	1
97	98	Female	36	Accountant	7.1	8	60	4	Normal	1
98	99	Female	36	Teacher	7.1	8	60	4	Normal	1
99	100	Female	36	Teacher	7.1	8	60	4	Normal	1

100 rows × 13 columns

In [73]:

```
b.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Person ID                            100 non-null    int64
1   Gender                               100 non-null    object
2   Age                                   100 non-null    int64
3   Occupation                           100 non-null    object
4   Sleep Duration                       100 non-null    float64
5   Quality of Sleep                     100 non-null    int64
6   Physical Activity Level              100 non-null    int64
7   Stress Level                         100 non-null    int64
8   BMI Category                         100 non-null    object
9   Blood Pressure                      100 non-null    object
10  Heart Rate                           100 non-null    int64
11  Daily Steps                          100 non-null    int64
12  Sleep Disorder                       100 non-null    object
dtypes: float64(1), int64(7), object(5)
memory usage: 10.3+ KB
```

In [74]:

```
b.describe()
```

Out[74]:

	Person ID	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	Heart Rate	
count	100.000000	100.000000	100.000000	100.000000	100.000000	100.000000	100.000000	
mean	50.500000	31.690000	6.871000	6.590000	51.910000	6.420000	71.610000	6
std	29.011492	2.26388	0.766903	1.005992	19.429279	1.485145	4.240009	4
min	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000	3
25%	25.750000	30.000000	6.100000	6.000000	30.000000	6.000000	70.000000	4
50%	50.500000	31.500000	7.100000	7.000000	60.000000	6.000000	70.000000	7
75%	75.250000	33.000000	7.700000	7.000000	75.000000	8.000000	72.000000	8
max	100.000000	36.000000	7.900000	8.000000	75.000000	8.000000	85.000000	10

In [75]:

```
b.columns
```

Out[75]:


```
Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',  
      'Quality of Sleep', 'Physical Activity Level', 'Stress Level',  
      'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps',  
      'Sleep Disorder'],  
      dtype='object')
```

In [76]:

```
sns.pairplot(b)
```

Out[76]:

<seaborn.axisgrid.PairGrid at 0x171dfb4b4f0>



In [77]:


```
sns.distplot(b['Heart Rate'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[77]:

```
<AxesSubplot:xlabel='Heart Rate', ylabel='Density'>
```



In [78]:


```
f=b[['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',  
    'Quality of Sleep', 'Physical Activity Level', 'Stress Level',  
    'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps',  
    'Sleep Disorder']]
```


In [79]:

```
sns.heatmap(f.corr())
```

Out[79]:

<AxesSubplot:>



In [80]:

```
x3=f[['Person ID','Age','Sleep Duration',
      'Quality of Sleep', 'Physical Activity Level', 'Stress Level',
      'Heart Rate', 'Daily Steps']]
y3=f['Heart Rate']
```

In [81]:

```
from sklearn.model_selection import train_test_split
x3_train,x3_test,y3_train,y3_test=train_test_split(x3,y3,test_size=0.5)
```

In [82]:

```
from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x3_train,y3_train)
```

Out[82]:

LinearRegression()

In [83]:

```
print(lr.intercept_)
```

2.8421709430404007e-13

In [84]:

```
r=pd.DataFrame(lr.coef_,x3.columns,columns=['Co-efficient'])  
r
```

Out[84]:


	Co-efficient
Person ID	3.338952e-16
Age	-3.645748e-15
Sleep Duration	-2.959317e-16
Quality of Sleep	-2.342765e-15
Physical Activity Level	2.160301e-17
Stress Level	-1.231670e-15
Heart Rate	1.000000e+00
Daily Steps	-2.874025e-17

In [85]:

```
u=lr.predict(x3_test)  
plt.scatter(y3_test,u)
```

Out[85]:

<matplotlib.collections.PathCollection at 0x171e3210eb0>



In [86]:

```
print(lr.score(x3_test,y3_test))
```

1.0

In [87]:

```
lr.score(x3_train,y3_train)
```

Out[87]:

1.0

In [88]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [89]:

```
rr=Ridge(alpha=10)  
rr.fit(x3_train,y3_train)
```

Out[89]:

Ridge(alpha=10)

In [90]:

```
rr.score(x3_test,y3_test)
```

Out[90]:

0.9995763294228194

In [91]:

```
la=Lasso(alpha=10)  
la.fit(x3_train,y3_train)
```

Out[91]:

Lasso(alpha=10)

In [92]:

```
la.score(x3_test,y3_test)
```

Out[92]:

0.28648433285844843

ELASTIC NET

In [94]:

```
from sklearn.linear_model import ElasticNet  
p=ElasticNet()  
p.fit(x3_train,y3_train)
```

Out[94]:

ElasticNet()

In [95]:

```
print(p.coef_)
```

```
[-8.11075063e-03 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00  
 0.00000000e+00 0.00000000e+00 8.87540114e-01 -2.15505154e-04]
```

In [96]:

```
print(p.intercept_)
```

9.881434969249142

In [97]:

```
print(p.predict(x3_test))
```

```
[70.04998994 67.94699457 72.17148786 70.1392082 72.21204161 72.0660481
 70.07432219 70.05810069 70.1554297 72.0984911 65.74782754 70.09054369
 72.41481038 70.2040942 69.88777493 72.22826311 69.95266093 67.93077307
 69.96077168 72.18770936 69.5309019 69.59578791 68.18220634 74.26755972
 65.78027055 69.93643943 69.38861301 79.98647584 69.84722118 70.00132544
 75.58956348 70.06621144 77.9664487 77.30879134 72.20393086 70.17976195
 72.40669962 68.02810208 65.76404904 84.64338618 70.1716512 70.04187919
 72.17959861 69.87966418 75.58145273 70.1229867 65.7883813 70.12111103
 72.10660185 72.1309341 ]
```

In [98]:

```
print(p.score(x3_test,y3_test))
```

0.9910400441354008

EVALUATION METRICS

In [99]:

```
from sklearn import metrics
```

In [100]:

```
print("Mean Absolytre Error:",metrics.mean_absolute_error(y3_test,prediction))
```

Mean Absolytre Error: 70.23635755897878

In [101]:

```
print("Mean Squared Error:",metrics.mean_squared_error(y3_test,prediction))
```

Mean Squared Error: 4947.82351088423

In [102]:

```
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y3_test,prediction)))
```

Root Mean Squared Error: 70.34076706209729

17_student_marks

In [103]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\17_student_marks.csv")  
a
```

Out[103]:

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6	Test_7	Test_8	Test_9	Test_10
0	22000	78	87	91	91	88	98	94	100	100	100
1	22001	79	71	81	72	73	68	59	69	59	60
2	22002	66	65	70	74	78	86	87	96	88	80
3	22003	60	58	54	61	54	57	64	62	72	60
4	22004	99	95	96	93	97	89	92	98	91	90
5	22005	41	36	35	28	35	36	27	26	19	20
6	22006	47	50	47	57	62	64	71	75	85	80
7	22007	84	74	70	68	58	59	56	56	64	70
8	22008	74	64	58	57	53	51	47	45	42	40
9	22009	87	81	73	74	71	63	53	45	39	40
10	22010	40	34	37	33	31	35	39	38	40	40
11	22011	91	84	78	74	76	80	80	73	75	70
12	22012	81	83	93	88	89	90	99	99	95	80
13	22013	52	50	42	38	33	30	28	22	12	20
14	22014	63	67	65	74	80	86	95	96	92	80
15	22015	76	82	88	94	85	76	70	60	50	50
16	22016	83	78	71	71	77	72	66	75	66	60
17	22017	55	45	43	38	43	35	44	37	45	30
18	22018	71	67	76	74	64	61	57	64	61	50
19	22019	62	61	53	49	54	59	68	74	65	50
20	22020	44	38	36	34	26	34	39	44	36	40
21	22021	50	56	53	46	41	38	47	39	44	30
22	22022	57	48	40	45	43	36	26	19	9	20
23	22023	59	56	52	44	50	40	45	46	54	50
24	22024	84	92	89	80	90	80	84	74	68	70
25	22025	74	80	86	87	90	100	95	87	85	70
26	22026	92	84	74	83	93	83	75	82	81	70
27	22027	63	70	74	65	64	55	61	58	48	40
28	22028	78	77	69	76	78	74	67	69	78	60
29	22029	55	58	59	67	71	62	53	61	67	70
30	22030	54	54	48	38	35	45	46	47	41	30
31	22031	84	93	97	89	86	95	100	100	100	90
32	22032	95	100	94	100	98	99	100	90	80	80
33	22033	64	61	63	73	63	68	64	58	50	50
34	22034	76	79	73	77	83	86	95	89	90	90
35	22035	78	71	61	55	54	48	41	32	41	40
36	22036	95	89	91	84	89	94	85	91	100	100

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6	Test_7	Test_8	Test_9	Test_10
37	22037	99	89	79	87	87	81	82	74	64	4
38	22038	82	83	85	86	89	80	88	95	87	9
39	22039	65	56	64	62	58	51	61	68	70	7
40	22040	100	93	92	86	84	76	82	74	79	7
41	22041	78	72	73	79	81	73	71	77	83	9
42	22042	98	100	100	93	94	92	100	100	98	9
43	22043	58	62	67	77	71	63	64	73	83	7
44	22044	96	92	94	100	99	95	98	92	84	8
45	22045	86	87	85	84	85	91	86	82	85	8
46	22046	48	55	46	40	34	29	37	34	39	4
47	22047	56	52	54	47	40	35	43	44	40	5
48	22048	42	44	46	53	62	59	57	53	43	5
49	22049	64	54	49	59	54	55	57	59	63	7
50	22050	50	44	37	29	37	46	53	57	55	6
51	22051	70	60	70	62	67	67	68	67	72	6
52	22052	63	73	70	63	60	67	61	59	52	5
53	22053	92	100	100	100	100	100	92	87	94	10
54	22054	64	55	54	61	63	57	47	37	44	4
55	22055	60	66	68	58	49	47	39	29	39	4

In [104]:

```
b=a.head(100)  
b
```


Out[104]:

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6	Test_7	Test_8	Test_9	Test_10
0	22000	78	87	91	91	88	98	94	100	100	100
1	22001	79	71	81	72	73	68	59	69	59	60
2	22002	66	65	70	74	78	86	87	96	88	80
3	22003	60	58	54	61	54	57	64	62	72	60
4	22004	99	95	96	93	97	89	92	98	91	90
5	22005	41	36	35	28	35	36	27	26	19	20
6	22006	47	50	47	57	62	64	71	75	85	80
7	22007	84	74	70	68	58	59	56	56	64	70
8	22008	74	64	58	57	53	51	47	45	42	40
9	22009	87	81	73	74	71	63	53	45	39	40
10	22010	40	34	37	33	31	35	39	38	40	40
11	22011	91	84	78	74	76	80	80	73	75	70
12	22012	81	83	93	88	89	90	99	99	95	80
13	22013	52	50	42	38	33	30	28	22	12	20
14	22014	63	67	65	74	80	86	95	96	92	80
15	22015	76	82	88	94	85	76	70	60	50	50
16	22016	83	78	71	71	77	72	66	75	66	60
17	22017	55	45	43	38	43	35	44	37	45	30
18	22018	71	67	76	74	64	61	57	64	61	50
19	22019	62	61	53	49	54	59	68	74	65	50
20	22020	44	38	36	34	26	34	39	44	36	40
21	22021	50	56	53	46	41	38	47	39	44	30
22	22022	57	48	40	45	43	36	26	19	9	10
23	22023	59	56	52	44	50	40	45	46	54	50
24	22024	84	92	89	80	90	80	84	74	68	70
25	22025	74	80	86	87	90	100	95	87	85	70
26	22026	92	84	74	83	93	83	75	82	81	70
27	22027	63	70	74	65	64	55	61	58	48	40
28	22028	78	77	69	76	78	74	67	69	78	60
29	22029	55	58	59	67	71	62	53	61	67	70
30	22030	54	54	48	38	35	45	46	47	41	30
31	22031	84	93	97	89	86	95	100	100	100	90
32	22032	95	100	94	100	98	99	100	90	80	80
33	22033	64	61	63	73	63	68	64	58	50	50
34	22034	76	79	73	77	83	86	95	89	90	90
35	22035	78	71	61	55	54	48	41	32	41	40
36	22036	95	89	91	84	89	94	85	91	100	100

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6	Test_7	Test_8	Test_9	Test_10
37	22037	99	89	79	87	87	81	82	74	64	81
38	22038	82	83	85	86	89	80	88	95	87	81
39	22039	65	56	64	62	58	51	61	68	70	71
40	22040	100	93	92	86	84	76	82	74	79	71
41	22041	78	72	73	79	81	73	71	77	83	81
42	22042	98	100	100	93	94	92	100	100	98	81
43	22043	58	62	67	77	71	63	64	73	83	71
44	22044	96	92	94	100	99	95	98	92	84	81
45	22045	86	87	85	84	85	91	86	82	85	81
46	22046	48	55	46	40	34	29	37	34	39	41
47	22047	56	52	54	47	40	35	43	44	40	41
48	22048	42	44	46	53	62	59	57	53	43	41
49	22049	64	54	49	59	54	55	57	59	63	71
50	22050	50	44	37	29	37	46	53	57	55	61
51	22051	70	60	70	62	67	67	68	67	72	61
52	22052	63	73	70	63	60	67	61	59	52	61
53	22053	92	100	100	100	100	100	92	87	94	100
54	22054	64	55	54	61	63	57	47	37	44	41
55	22055	60	66	68	58	49	47	39	29	39	41

In [105]:

b.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56 entries, 0 to 55
Data columns (total 13 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Student_ID  56 non-null    int64
 1   Test_1      56 non-null    int64
 2   Test_2      56 non-null    int64
 3   Test_3      56 non-null    int64
 4   Test_4      56 non-null    int64
 5   Test_5      56 non-null    int64
 6   Test_6      56 non-null    int64
 7   Test_7      56 non-null    int64
 8   Test_8      56 non-null    int64
 9   Test_9      56 non-null    int64
10  Test_10     56 non-null    int64
11  Test_11     56 non-null    int64
12  Test_12     56 non-null    int64
dtypes: int64(13)
memory usage: 5.8 KB

```

In [106]:

```
b.describe()
```

Out[106]:

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6
count	56.000000	56.000000	56.000000	56.000000	56.000000	56.000000	56.000000
mean	22027.500000	70.750000	69.196429	68.089286	67.446429	67.303571	66.000000
std	16.309506	17.009356	17.712266	18.838333	19.807179	20.746890	21.054043
min	22000.000000	40.000000	34.000000	35.000000	28.000000	26.000000	29.000000
25%	22013.750000	57.750000	55.750000	53.000000	54.500000	53.750000	50.250000
50%	22027.500000	70.500000	68.500000	70.000000	71.500000	69.000000	65.500000
75%	22041.250000	84.000000	83.250000	85.000000	84.000000	85.250000	83.750000
max	22055.000000	100.000000	100.000000	100.000000	100.000000	100.000000	100.000000

In [107]:

```
b.columns
```

Out[107]:


```
Index(['Student_ID', 'Test_1', 'Test_2', 'Test_3', 'Test_4', 'Test_5',  
      'Test_6', 'Test_7', 'Test_8', 'Test_9', 'Test_10', 'Test_11',  
      'Test_12'],  
      dtype='object')
```

In [108]:

```
sns.pairplot(b)
```

Out[108]:

<seaborn.axisgrid.PairGrid at 0x171e3238580>



In [109]:


```
sns.distplot(b['Test_2'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[109]:

<AxesSubplot:xlabel='Test_2', ylabel='Density'>



In [110]:


```
f=b[['Student_ID', 'Test_1', 'Test_2', 'Test_3', 'Test_4', 'Test_5',  
     'Test_6', 'Test_7', 'Test_8', 'Test_9', 'Test_10', 'Test_11',  
     'Test_12']]
```

In [111]:

```
sns.heatmap(f.corr())
```

Out[111]:

<AxesSubplot:>



In [112]:

```
x4=f[['Student_ID', 'Test_1','Test_3', 'Test_4', 'Test_5',
      'Test_6', 'Test_7', 'Test_8', 'Test_9', 'Test_10', 'Test_11',
      'Test_12']]
y4=f['Test_2']
```

In [113]:

```
from sklearn.model_selection import train_test_split
x4_train,x4_test,y4_train,y4_test=train_test_split(x4,y4,test_size=0.5)
```

In [114]:

```
from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()
lr.fit(x4_train,y4_train)
```

Out[114]:

LinearRegression()

In [115]:

```
print(lr.intercept_)
```

-1936.3869042265571

In [116]:

```
r=pd.DataFrame(lr.coef_,x4.columns,columns=['Co-efficient'])  
r
```

Out[116]:


	Co-efficient
Student_ID	0.088189
Test_1	0.371436
Test_3	0.492043
Test_4	-0.063353
Test_5	0.261007
Test_6	-0.180162
Test_7	0.210883
Test_8	-0.176475
Test_9	0.069315
Test_10	0.058202
Test_11	0.067368
Test_12	-0.191963

In [118]:

```
u=lr.predict(x4_test)  
plt.scatter(y4_test,u)
```

Out[118]:

<matplotlib.collections.PathCollection at 0x171ec20b070>



In [119]:

```
print(lr.score(x4_test,y4_test))
```

0.9075725286491936

In [120]:

```
lr.score(x4_train,y4_train)
```

Out[120]:

0.9860267918987687

In [121]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [122]:

```
rr=Ridge(alpha=10)  
rr.fit(x4_train,y4_train)
```

Out[122]:

Ridge(alpha=10)

In [124]:

```
rr.score(x4_test,y4_test)
```

Out[124]:

0.9081261200925351

In [125]:

```
la=Lasso(alpha=10)  
la.fit(x4_train,y4_train)
```

Out[125]:

Lasso(alpha=10)

In [126]:

```
la.score(x4_test,y4_test)
```

Out[126]:

0.9103242376616466

ELASTIC NET

In [170]:

```
from sklearn.linear_model import ElasticNet  
p=ElasticNet()  
p.fit(x4_train,y4_train)
```

Out[170]:

ElasticNet()

In [171]:

```
print(p.coef_)
```

```
[ 0.07928517  0.35948343  0.51489021 -0.          0.15815984 -0.10461042
 0.09059163 -0.07292736  0.05057222  0.          0.07337492 -0.14493778]
```

In [172]:

```
print(p.intercept_)
```

```
-1740.5648305729223
```

In [173]:

```
print(p.predict(x4_test))
```

```
[49.88745773 70.29717482 55.08295575 86.1251413  73.98719324 75.54328634
 65.1077985  89.9578371  72.44900834 57.96129177 74.95397091 47.17348935
 52.43330373 37.939942   62.72285822 89.63881295 56.77317452 94.38943327
 64.61147947 82.40823394 63.55055548 70.91961992 64.89729884 57.696127
 72.7367411  76.38735574 66.44555203 95.03758024]
```

In [174]:

```
prediction=p.predict(x4_test)
print(p.score(x4_test,y4_test))
```

```
0.9100721717195225
```

EVALUATION METRICS

In [175]:

```
from sklearn import metrics
```

In [176]:

```
print("Mean Absolytre Error:",metrics.mean_absolute_error(y4_test,prediction))
```

```
Mean Absolytre Error: 3.9116681158081894
```

In [177]:

```
print("Mean Squared Error:",metrics.mean_squared_error(y4_test,prediction))
```

```
Mean Squared Error: 22.380558855578638
```

In [178]:

```
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y4_test,prediction)))
```

```
Root Mean Squared Error: 4.730809534908231
```

18_world-data-2023


In [136]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\18_world-data-2023.csv")
a
```

Out[136]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Call Co
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	9
1	Albania	105	AL	43.10%	28,748	9,000	11.78	35
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	21
3	Andorra	164	AD	40.00%	468	NaN	7.20	37
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	24
...
190	Venezuela	32	VE	24.50%	912,050	343,000	17.88	5
191	Vietnam	314	VN	39.30%	331,210	522,000	16.75	8
192	Yemen	56	YE	44.60%	527,968	40,000	30.45	96
193	Zambia	25	ZM	32.10%	752,618	16,000	36.19	26
194	Zimbabwe	38	ZW	41.90%	390,757	51,000	30.68	26

195 rows × 35 columns




In [137]:

```
b=a.head(100)
b
```

Out[137]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Call Co
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	9
1	Albania	105	AL	43.10%	28,748	9,000	11.78	35
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	21
3	Andorra	164	AD	40.00%	468	NaN	7.20	37
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	24
...
95	Lesotho	71	LS	77.60%	30,355	2,000	26.81	26
96	Liberia	53	LR	28.00%	111,369	2,000	33.04	23
97	Libya	4	LY	8.70%	1,759,540	0	18.83	21
98	Liechtenstein	238	LI	32.20%	160	NaN	9.90	42
99	Lithuania	43	LT	47.20%	65,300	34,000	10.00	37

100 rows × 35 columns



In [138]:

b.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 35 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Country                             100 non-null    object
 1   Density (P/Km2)                     100 non-null    object
 2   Abbreviation                        96 non-null     object
 3   Agricultural Land( %)              98 non-null     object
 4   Land Area(Km2)                     100 non-null    object
 5   Armed Forces size                   92 non-null     object
 6   Birth Rate                          98 non-null     float64
 7   Calling Code                       100 non-null    float64
 8   Capital/Major City                 99 non-null     object
 9   Co2-Emissions                      98 non-null     object
10   CPI                                93 non-null     object
11   CPI Change (%)                     94 non-null     object
12   Currency-Code                      91 non-null     object
13   Fertility Rate                     98 non-null     float64
14   Forested Area (%)                  98 non-null     object
15   Gasoline Price                     93 non-null     object
16   GDP                                99 non-null     object
17   Gross primary education enrollment (%) 97 non-null     object
18   Gross tertiary education enrollment (%) 95 non-null     object
19   Infant mortality                   97 non-null     float64
20   Largest city                       97 non-null     object
21   Life expectancy                    97 non-null     float64
22   Maternal mortality ratio           95 non-null     float64
23   Minimum wage                       80 non-null     object
24   Official language                  100 non-null    object
25   Out of pocket health expenditure    97 non-null     object
26   Physicians per thousand             97 non-null     float64
27   Population                         100 non-null    object
28   Population: Labor force participation (%) 92 non-null     object
29   Tax revenue (%)                    87 non-null     object
30   Total tax rate                     96 non-null     object
31   Unemployment rate                  92 non-null     object
32   Urban_population                   98 non-null     object
33   Latitude                           100 non-null    float64
34   Longitude                           100 non-null    float64
dtypes: float64(9), object(26)
memory usage: 27.5+ KB

```

In [139]:

```
b.describe()
```

Out[139]:

	Birth Rate	Calling Code	Fertility Rate	Infant mortality	Life expectancy	Maternal mortality ratio	Physicians per thousand	
count	98.000000	100.00000	98.000000	97.000000	97.000000	95.000000	97.000000	100
mean	20.045000	354.15000	2.651633	21.782474	72.358763	164.305263	1.916186	20
std	9.737986	330.66234	1.239758	20.437751	7.712838	226.488977	1.809687	20
min	7.200000	1.00000	1.270000	1.400000	52.800000	2.000000	0.040000	-30
25%	11.420000	84.75000	1.695000	5.000000	66.600000	12.000000	0.280000	0
50%	18.125000	253.50000	2.180000	12.800000	74.100000	58.000000	1.580000	10
75%	27.027500	501.25000	3.322500	32.900000	78.100000	244.500000	3.190000	40
max	42.170000	1876.00000	5.920000	84.500000	84.200000	1140.000000	8.420000	60

In [140]:

```
c=b.dropna(axis=1)
c
```

Out[140]:

	Country	Density\n(P/Km2)	Land Area(Km2)	Calling Code	Official language	Population	Latitude	Longitude
0	Afghanistan	60	652,230	93.0	Pashto	38,041,754	33.939110	67.159810
1	Albania	105	28,748	355.0	Albanian	2,854,191	41.153332	20.165321
2	Algeria	18	2,381,741	213.0	Arabic	43,053,054	28.033886	0.839269
3	Andorra	164	468	376.0	Catalan	77,142	42.506285	1.521106
4	Angola	26	1,246,700	244.0	Portuguese	31,825,295	-11.202692	17.084497
...
95	Lesotho	71	30,355	266.0	English	2,125,268	-29.609988	28.249686
96	Liberia	53	111,369	231.0	English	4,937,374	6.428055	-9.432580
97	Libya	4	1,759,540	218.0	Arabic	6,777,452	26.335100	17.054833
98	Liechtenstein	238	160	423.0	German	38,019	47.141039	9.683156
99	Lithuania	43	65,300	370.0	Lithuanian	2,786,844	55.169438	25.264604

100 rows × 8 columns

In [141]:

```
c.columns
```

Out[141]:


```
Index(['Country', 'Density\n(P/Km2)', 'Land Area(Km2)', 'Calling Code',  
      'Official language', 'Population', 'Latitude', 'Longitude'],  
      dtype='object')
```

In [142]:

```
sns.pairplot(c)
```

Out[142]:

<seaborn.axisgrid.PairGrid at 0x171ec34a2e0>



In [143]:


```
sns.distplot(c['Calling Code'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[143]:

```
<AxesSubplot:xlabel='Calling Code', ylabel='Density'>
```



In [144]:


```
f=c[['Country', 'Density\n(P/Km2)', 'Land Area(Km2)', 'Calling Code',  
     'Official language', 'Population', 'Latitude', 'Longitude']]
```

In [145]:

```
sns.heatmap(f.corr())
```

Out[145]:

```
<AxesSubplot:>
```



In [146]:

```
x=f[['Latitude', 'Longitude']]
y=f['Calling Code']
```

In [147]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.5)
```

In [148]:

```
from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[148]:

LinearRegression()

In [149]:

```
print(lr.intercept_)
```

362.5847124562507

In [150]:

```
r=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
r
```

Out[150]:


	Co-efficient
Latitude	-0.498721
Longitude	1.511282

In [151]:

```
u=lr.predict(x_test)
plt.scatter(y_test,u)
```

Out[151]:

<matplotlib.collections.PathCollection at 0x171eca85520>



In [152]:

```
print(lr.score(x_test,y_test))
```

-0.10366393692391407

In [153]:

```
lr.score(x_train,y_train)
```

Out[153]:

0.06275371502200777

In [154]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [155]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[155]:

Ridge(alpha=10)

In [156]:

```
rr.score(x_test,y_test)
```

Out[156]:

-0.10364803763482766

In [157]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[157]:

Lasso(alpha=10)

In [158]:

```
la.score(x_test,y_test)
```

Out[158]:

-0.10279538065354266

ELASTIC NET

In [159]:

```
from sklearn.linear_model import ElasticNet
p=ElasticNet()
p.fit(x_train,y_train)
```

Out[159]:

ElasticNet()

In [160]:

```
print(p.coef_)
```

[-0.49725868 1.51090022]

In [161]:

```
print(p.intercept_)
```

362.5559749931516

In [162]:

```
print(p.predict(x_test))
```

```
[395.75872307 377.5580689 397.43924134 386.58337386 377.24886466
419.72070815 368.57156158 379.8127267 343.71864512 301.51395
371.20471649 357.0585514 427.54819658 344.1961059 410.95537202
231.16516868 553.43918496 221.3829101 533.6428293 363.09529168
378.08308861 399.90575524 447.98244062 415.11503162 407.02220901
352.90460375 332.73914364 535.07220671 485.51046653 439.79283411
123.86187229 285.54026919 455.03085304 319.51805242 502.15230125
243.90664037 263.34012556 421.02547825 248.02721451 419.18440012
380.49632781 371.4450462 414.47845684 271.10146908 402.09986423
345.11253672 433.81909344 577.24477111 367.42973912 370.66656053]
```

In [163]:

```
prediction=p.predict(x_test)
print(p.score(x_test,y_test))
```

-0.10358069305318973

EVALUATION METRICS

In [164]:

```
from sklearn import metrics
```

In [165]:

```
print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolytre Error: 255.90098769716585

In [166]:

```
print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

Mean Squared Error: 99755.58333391439

In [167]:

```
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root Mean Squared Error: 315.84107290521035

In []: