

DATA COLLECTION

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\11_winequality-red.csv")
a
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
...	
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	

1599 rows × 12 columns



In [3]:

```
a.head(5)
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alco
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	

DATA CLEANING AND PRE-PROCESSING

In [4]:

```
a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide     1599 non-null   float64
6   total sulfur dioxide    1599 non-null   float64
7   density                1599 non-null   float64
8   pH                     1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

In [5]:

```
a.describe()
```

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.406827
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.847047
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000

In [6]:

```
a.columns
```

Out[6]:

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
      'pH', 'sulphates', 'alcohol', 'quality'],  
      dtype='object')
```

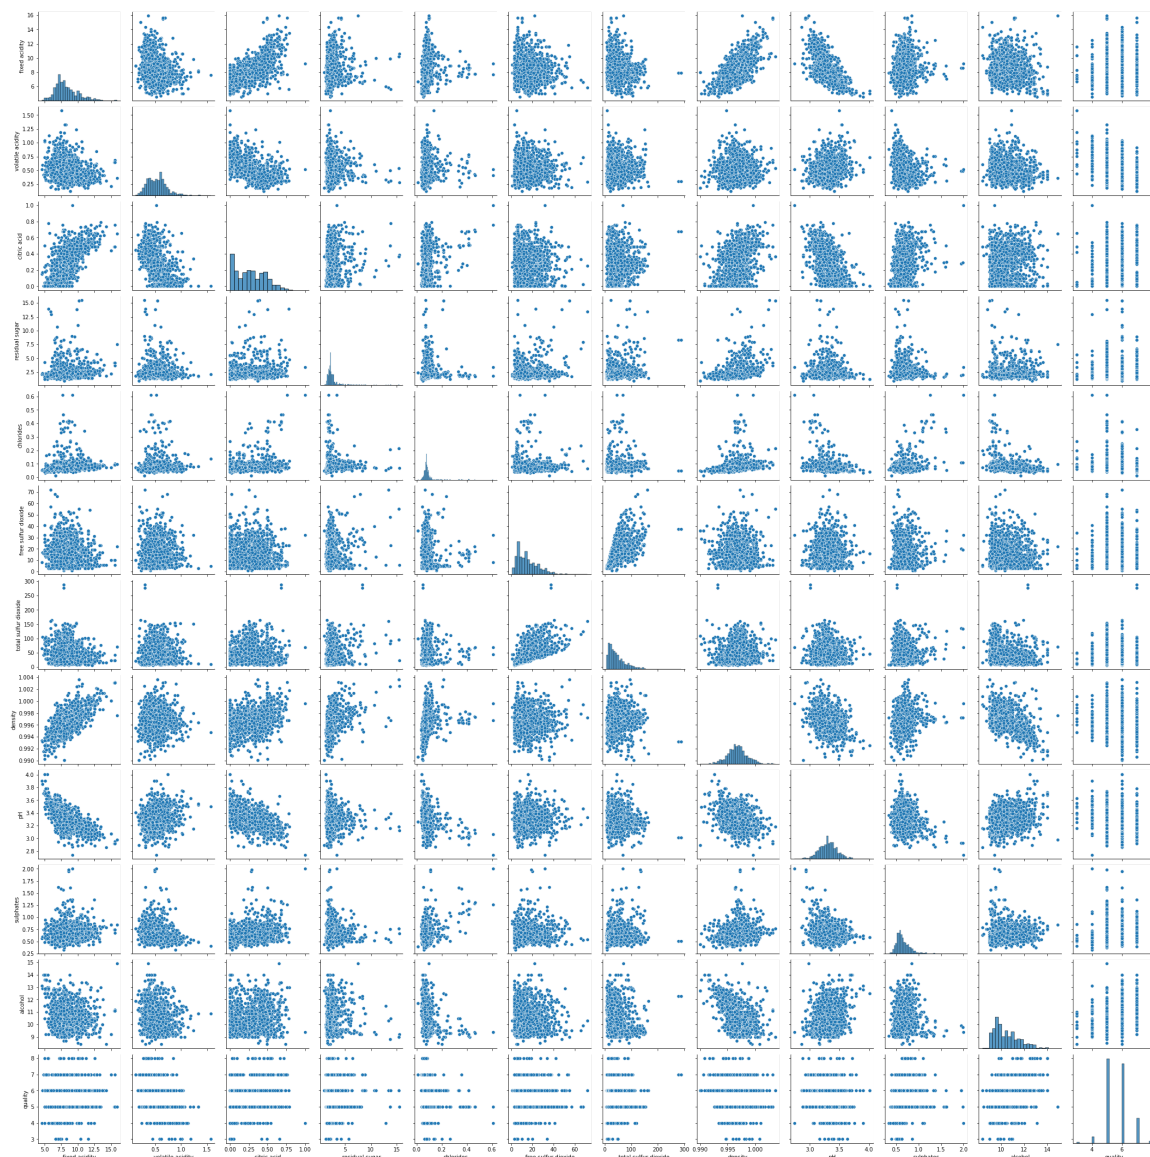
EDA and VISUALIZATION

In [7]:

```
sns.pairplot(a)
```

Out[7]:

<seaborn.axisgrid.PairGrid at 0x18d4f636d90>



In [8]:


```
sns.distplot(a['citric acid'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

warnings.warn(msg, FutureWarning)

Out[8]:

<AxesSubplot:xlabel='citric acid', ylabel='Density'>



In [9]:


```
f=a[['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
    'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
    'pH', 'sulphates', 'alcohol', 'quality']]
```

In [10]:

```
sns.heatmap(f.corr())
```

Out[10]:

<AxesSubplot:>



To train the model-model building

In [11]:

```
x=f[['fixed acidity', 'volatile acidity','residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
      'pH', 'sulphates', 'alcohol', 'quality']]
y=f['citric acid']
```

In [12]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [13]:

```
from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[13]:

LinearRegression()

In [14]:

```
print(lr.intercept_)
```

2.5843290594097157

In [15]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[15]:


	Co-efficient
fixed acidity	0.061519
volatile acidity	-0.446338
residual sugar	0.005519
chlorides	0.825668
free sulfur dioxide	-0.002472
total sulfur dioxide	0.001272
density	-2.911242
pH	-0.013626
sulphates	0.027735
alcohol	0.022281
quality	-0.000503

In [16]:

```
prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[16]:

<matplotlib.collections.PathCollection at 0x18d58ebb160>



In [17]:

```
print(lr.score(x_test,y_test))
```

0.6876992302836324

In [18]:

```
lr.score(x_train,y_train)
```

Out[18]:

0.6743276120005179

RIDGE REGRESSION

In [19]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [20]:

```
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[20]:

Ridge(alpha=10)

In [21]:

```
rr.score(x_test,y_test)
```

Out[21]:

0.6775140538038511

LASSO REGRESSION

In [22]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[22]:

Lasso(alpha=10)

In [23]:

```
la.score(x_test,y_test)
```

Out[23]:

-0.007117544613139071

In [24]:

```
from sklearn.linear_model import ElasticNet
p=ElasticNet()
p.fit(x_train,y_train)
```

Out[24]:

ElasticNet()

In [25]:

```
print(p.coef_)
[ 0. -0.  0.  0. -0.  0.  0. -0.  0.  0.  0.]
```

In [26]:

```
print(p.intercept_)
0.2659338695263629
```

In [27]:

```
print(p.predict(x_test))
```

localhost:8888/notebooks/11.winequality-red ee.ipynb

In [28]:

-0.007117544613139071

In [29]:

In [30]:

Mean Absolytre Error: 0.1686069221030682

In [31]:

Mean Squared Error: 0.03991400865533665

In [32]:

Root Mean Squared Error: 0.19978490597474236

In []: